



FAST-TIME SIMULATION EVALUATION OF A CONFLICT RESOLUTION ALGORITHM UNDER HIGH AIR TRAFFIC DEMAND

Todd Farley, NASA Ames Research Center, Moffett Field, California

Heinz Erzberger, Adjunct Professor, University of California, Santa Cruz, California

Abstract

In this paper, an automated conflict resolution algorithm is evaluated based on fast-time simulations of nominal and heavily increased air traffic demand in the Cleveland Air Route Traffic Control Center airspace. The algorithm under study is designed to support an automated separation assurance capability for next-generation air traffic management systems. It resolves detected conflicts that are projected to be between one and twenty minutes from first loss of separation. Rule bases are used to determine which aircraft to maneuver and which types of maneuver to consider: climb/descent, path stretch, or speed change. The algorithm uses high-fidelity trajectory modeling to identify a four-dimensional resolution trajectory that begins at the aircraft's current position and altitude, is conflict-free for a specified period of time, and ends at a position and altitude on the aircraft's original trajectory. Two case studies are presented to illustrate the operation of the subject algorithm: a typical vertical conflict involving traffic descending through busy flight levels, and an arrival conflict with arrival-fix crossing restrictions and sequencing constraints. The simulation environment is a medium-fidelity, fast-time simulation of departure, en-route, and arrival traffic based on recorded FAA data, and it assumes that all flights adhere to their four-dimensional trajectories precisely. Within the limitations of the simulation, the results indicate that the conflict resolution algorithm is capable of resolving conflicts safely and efficiently at traffic levels significantly higher than today. Safety and efficiency metrics are offered as benchmarks for comparison with alternative algorithms.

Introduction

Conventional methods for assuring safe separation¹ of aircraft from proximate traffic are not scalable to accommodate the level of air traffic

demand that is projected for the years 2020 and beyond [1]. To meet the demand forecasts, air traffic management (ATM) system designers in Europe and the US are developing plans for next-generation ATM systems that rely on a separation assurance capability that is more scalable than that of today's system. This future capability is expected to rely on advanced automation to detect conflicts and to resolve them [2]. The design standards for an *automated separation assurance* function are calling for safe and efficient separation of air traffic at two- to three-times recent levels in order to adequately relieve the constraints that conventional, human-based separation assurance methods are expected to present as traffic levels increase. Accordingly, the international research community is investigating conflict detection and conflict resolution algorithms that will meet the two- to three-fold traffic-increase design standard. In this paper, we focus specifically on the conflict resolution problem.

Many candidate algorithms for conflict resolution have been proposed. Kuchar and Yang cite 47 of them in a broad survey published in 2000 [3].² Many of the cited investigations remain active as of this writing, and the literature has expanded to include approaches based on genetic algorithms (for example, [4] and [5]) and the incorporation of arrival-time constraints (such as [6] and [7]).

Most proposed solutions have been exercised only through a set of scripted, constrained or simplified examples. Such exercises are a necessary step in the validation process. However, in order to validate a conflict resolution algorithm with respect to the "2x-to-3x" design standard, a candidate algorithm must be stressed by traffic volumes, densities, and complexities that are commensurate

¹ Two aircraft are said to have *safe separation* if they are separated vertically by at least 1000 ft (305 m) or they are separated horizontally by at least 5 nmi (9.3 km).

² To aid in comparisons, we adopt the Kuchar & Yang framework for categorization of conflict detection and resolution algorithms. The framework categorizes conflict resolution algorithms along three axes: scope (i.e., pairwise vs. global); strategy (i.e., prescribed, optimized, force field, or manual); and available maneuvers (e.g., turns, vertical maneuvers, speed changes, or combined maneuvers).

with today's busiest airspace as well as that of the envisioned future. Further, it is necessary to expose the algorithm to the full breadth and variety of conflict situations that occur—or, it is reasoned, may occur—in real-world operations now or in the foreseeable future. As of this writing, we are aware of no conflict resolution algorithm that has been evaluated under such conditions, and that provides the motivation for the work presented here.

Objectives

The present study conducts such an evaluation for one candidate algorithm: the conflict resolution algorithm developed as part of the Automated Airspace Concept [8][9]. Its results provide objective benchmarks by which ATM researchers may evaluate competing algorithms and/or approaches to the conflict-resolution problem. It is hoped that this work may stimulate the introduction of additional benchmarks generic enough to facilitate performance comparisons between alternative approaches. In so doing, it will help determine the potential of the present algorithm for further development.

Organization

The body of this report is organized into the following sections:

- An **overview** of the conflict resolution algorithm being evaluated in this study is provided in order to familiarize the reader with the design approach and its relative strengths and frailties; two case studies are included to put the approach in context;
- The experimental **approach** and **procedure** are presented to identify the underlying simplifications and assumptions made for this simulation-based study;
- The **results** are presented, categorized according to safety- and efficiency-related metrics, and discussed to help put the results into context;
- We close with a **summary** of the study's findings and suggestions for future research.

Overview of the Conflict Resolution Algorithm

The conflict resolution algorithm evaluated in this study represents the continued evolution of the algorithm developed by Erzberger (most recently, [10]). According to the Yang & Kuchar taxonomy, the algorithm may be categorized as a “pairwise,” “optimized” approach, with resolution trajectories comprised of turn maneuvers, vertical maneuvers, or speed changes. It employs an iterative approach, sequentially computing and evaluating candidate trajectories (we call them *trial-plan trajectories*) in a structured progression—where the progression is tailored according to the dynamic and physical characteristics of the conflict—until a trajectory is found that satisfies all of the resolution conditions. That trial-plan trajectory is then output by the algorithm as the *conflict resolution trajectory*. If no acceptable resolution is found, the algorithm employs one of three fault-recovery strategies: deferral, de-scope, or delegate. Figure 1 presents a flow chart of the algorithm. The remainder of this overview section explains and elaborates on the elements and flow of the algorithm as depicted in Figure 1.

The algorithm is invoked when a conflict detection algorithm (a separate element of the larger separation assurance system) identifies a pair of aircraft that is projected to violate the separation standard within a time interval of interest³. The primary inputs to the resolution algorithm are the conflict data that correspond to the aircraft projected to be in conflict. The conflict data include the flight identifiers, the predicted time until first loss of separation (t_{los}), state information for each aircraft (at present (t_0) and at the time when safe separation is predicted to be lost (t_0+t_{los})), and intent information (e.g., flight plans). The conflict data also include these same parameters for any additional conflicts—referred to as *downstream conflicts* (ref. Figure 2)—that are predicted to occur downstream of the *primary conflict*. Downstream conflicts must also be resolved if they occur within a specified time range⁴. Downstream conflicts, which occur with increased frequency in dense traffic, add complexity to the resolution process.

³ The simulation used a simple, deterministic conflict detection algorithm.

⁴ See *conflict-free time horizon* (CFH) on the next page.

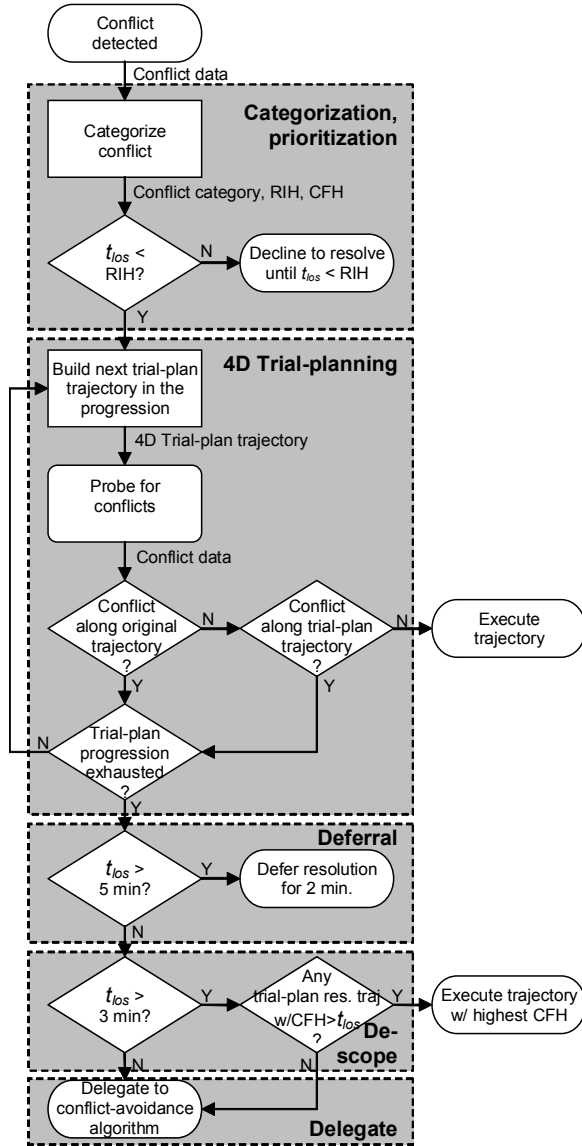


Figure 1. Flow chart of conflict resolution algorithm

Categorization and Prioritization

Based upon the physical and dynamic characteristics of the projected conflict as captured in the conflict data, the conflict is categorized and prioritized. Some example factors considered in the categorization include phase of flight, geometry of the conflict, and the number of prior trajectory amendments incurred by the aircraft. Categorization is used to select appropriate parameter values (explained below) and to inform the selection of which aircraft to maneuver and how. Prioritization is primarily a function of the time until first loss of separation (t_{los}), and t_{los} can also influence the resolution strategy as described later.

Among the parameters determined based on the category of the conflict are the *resolution-initiation* time horizon (RIH) and the *conflict-free* time horizon (CFH). The resolution-initiation time horizon is defined as the maximum time before initial loss of separation (as projected by the conflict detection function) that a detected conflict shall be made eligible for resolution. A larger resolution-initiation time horizon will invoke the conflict-resolution algorithm earlier, expanding the set of possible resolution maneuvers but also increasing its exposure to false conflict alerts and, therefore, unnecessary trajectory amendments and/or delay. Conversely, a smaller resolution-initiation time horizon reduces the exposure to false alerts but also decreases the time (and thus options) available to resolve genuine conflicts. Performance studies [11] led us to select a resolution-initiation time horizon of 8 minutes for most conflict categories. Twenty minutes is used for conflicts between aircraft on descent to the same arrival fix, where early action can significantly improve the likelihood of finding a resolution trajectory. Sensitivity studies are planned in order to improve our understanding of the tradeoffs associated with the selection of this parameter.

The conflict-free time horizon is defined as the future time period for which the algorithm shall guarantee a resolution trajectory to be conflict-free. Like a rain-free weather forecast, confidence in the conflict-free guarantee naturally diminishes as the conflict-free time horizon increases, owing to the accumulation of uncertainty over the longer time horizon. Resolution trajectories with longer conflict-free time horizons also are more difficult to find, because the likelihood that multiple conflicts will have to be resolved is increased. In simulations of the resolution algorithm described later, a CFH value of 12 min resulted in acceptable performance. Twenty minutes is used for conflicts between aircraft bound for the same arrival fix. As with the resolution-initiation time horizon, sensitivity studies are planned in order to improve our understanding of the tradeoffs associated with the selection of the conflict-free time horizon.

4D Trial-Planning Loop

The core of the conflict resolution algorithm is the trial-planning loop. This loop is explained in detail by Erzberger [10]. To summarize, the algorithm uses the results of the categorization process described above as well as trial-planning feedback to navigate a rule-based decision tree that guides selection of the *maneuver aircraft* and a *trial-plan maneuver*. Trial-plan maneuvers alter the

vertical, lateral, or speed profile of the maneuvering aircraft and return it to a point on its original trajectory, downstream of the projected conflict. Trial-plan maneuvers are prioritized on the basis of delay and deviation from the nominal flight plan trajectory, among other considerations. The progression of prioritized aircraft–maneuver pairs serves as a reservoir of alternative solutions through which to iterate in the course of attempting to resolve a conflict.




Figure 2. Conflict types, illustrated

Given a specific trial-plan maneuver for a specific aircraft, a 4D-trajectory synthesizer uses detailed models of aircraft performance, operational procedures, and the atmosphere (including winds aloft) to generate a flyable four-dimensional trajectory. Each trial-plan trajectory is checked for conflicts by comparing it against the 4D trajectories of all other aircraft in the airspace of interest. If the trial-plan trajectory is projected to be free of conflicts for the specified conflict-free time horizon, then the algorithm promotes the trial-plan trajectory to the status of *resolution trajectory*. The resolution trajectory is ready for execution by the maneuver aircraft.

If the trial-plan trajectory is projected to have a conflict within the specified conflict-free time horizon, appropriate diagnostic information is logged, which will retire that trial-plan trajectory and inform the selection of a new aircraft–maneuver pair from its reservoir. This iterative process continues until either (1) an acceptable resolution trajectory is found, or (2) the reservoir of available trial-plan maneuvers and aircraft combinations is exhausted.

If no resolution is found, the algorithm immediately transitions to fault-recovery mode, sequentially invoking one of three backup strategies in the following order: (1) defer, (2) de-scope, (3) delegate. The paragraphs that follow describe how RIH and CFH are used to manage this fault-recovery process.

Defer

The algorithm begins the fault-recovery sequence by adopting a *deferral* strategy. Quite simply, it waits a specified amount of time (we used 2 min in this study) and tries again from scratch. This process may be thought of as a temporary reduction of the resolution-initiation time horizon for this conflict pair.

The rationale behind the deferral strategy is that the marginal benefit of updated, more certain state or intent information will more than compensate for the marginal risk assumed by deferring action. New maneuvering lanes may become available, or the primary conflict itself may be resolved without action. The deferral process is repeated (i.e., the RIH is decremented by 2 min) until a resolution is found or until the time to first loss (t_{los}) has decreased to less than 5 min. The latter two values (the decrement and the t_{los} threshold) deserve further analysis to determine optimal values.

De-Scope

The deferral strategy above is a bet; if the primary conflict persists (i.e., the bet is lost), then t_{los} is now less than 5 min and the algorithm moves to a satisficing approach we call *de-scoping*. This strategy attempts to find a resolution by relaxing some longer-term constraints on the solution space, thus “de-scoping” the problem to a shorter time horizon. Specifically, the conflict-free time horizon is iteratively decremented until a resolution is found or the conflict-free time horizon has decreased to 3 min.

Delegate

In the event that neither deferral nor de-scoping is successful, then the conflict resolution algorithm has failed, and the conflict is delegated to a separate *conflict avoidance* algorithm. The present work relies on a conflict avoidance algorithm called the Tactical Separation Assisted Flight Environment, or TSAFE (described in the references on the design of Automated Airspace Concept [8][9]). Conflict avoidance algorithms fill the niche between conflict resolution (from 20 min to 3 min prior to loss of standard separation) and *collision avoidance*, such as TCAS (Traffic Alert and Collision Avoidance System), designed to avoid imminent collisions as

opposed to separation violations. Conflict avoidance algorithms sacrifice efficiency for expediency. They make greater use of the aircraft's performance envelope, and they waive certain constraints on the resolution, such as the need to return to the original flight plan.

Case Studies

Two case studies are presented to illustrate the operation of the subject algorithm: a typical conflict in the vertical dimension involving traffic descending through busy flight levels, and a more constrained example involving two arrival aircraft with arrival-fix crossing restrictions and sequencing constraints. The aircraft identifiers (call signs) have been removed; the identity of the specific operators is not germane to the discussion.

Case Study #1: Overflight vs. Arrival Conflict

Consider the conflict presented in Figure 3. This is a typical case of an arrival aircraft needing to descend through a stream of traffic at lower altitude on the same airway. Flight 1 (the **red** trajectory in the figure), an Airbus 320, is in cruise westbound at 34,000 feet on jet route J60 just north of Pittsburgh, Pennsylvania en route to Long Beach, California. Flight 2, a Hawker 800 twinjet (**blue** trajectory), is also in cruise westbound on J60 north of Pittsburgh, but at 38,000 feet. The flights are initially in trail but at different flight levels, and Flight 2 is about to begin its descent for arrival to Toledo, Ohio. The conflict detection algorithm projects that the descent trajectory of Flight 2 will conflict with the cruise trajectory of Flight 1, with first loss of separation occurring 6 min, 50 sec from now, when the aircraft positions and altitudes coincide with the points indicated with asterisks in the figures. The closest point of approach is predicted to be 1.1 nmi laterally and 875 ft vertically.

During categorization (ref. Figure 1), the conflict resolution algorithm classified this conflict as an “overflight-versus-arrival” conflict. Because aircraft on arrival are considered to be in a more critical stage of flight than aircraft in cruise, the algorithm is predisposed toward maneuvering the cruise aircraft, Flight 1, as opposed to the arrival aircraft, Flight 2.

Table 1 lists the progression of aircraft–maneuver pairings that the algorithm generated and evaluated during its trial-planning loop. The progression is a function of the conflict category, but it may also be altered based on feedback from prior trial plans in the progression.




Figure 3. Horizontal and vertical flight-plan trajectories of Flight 1 and Flight 2

Table 1. Progression of aircraft/maneuver pairings

Step	Aircraft	Maneuver	Result
1	Flight 1	Step climbs	Insufficient horiz. separation
2	Flight 1	Step descends	Secondary conflict
3	Flight 2	Temporary altitudes	Insufficient horiz. separation
4	Flight 1	40-sec path stretches (L/R)	Successful
5	Flight 1	Up to 3-min path stretches (L/R)	<not needed>
6	Flight 2	40-sec path stretches (L/R)	<not needed>
7	Flight 2	Up to 3-min path stretches (L/R)	<not needed>

In this case, the algorithm began by computing step-climb trajectories of 1000 and 2000 ft for Flight 1. Neither trajectory resolved the conflict because the two aircraft would have remained too close in trail. The algorithm then progressed to step-descent trajectories. A 1000-ft step descent was projected to successfully resolve the conflict with Flight 2 due to the associated increase in airspeed; however, it also was projected to cause a conflict 4 minutes later with a third aircraft, Flight 3, en route from Detroit to Harrisburg as shown in Figure 4. Therefore, this trial was rejected. Larger step-descent trajectories were unsuccessful in establishing safe separation between Flight 1 and Flight 2.

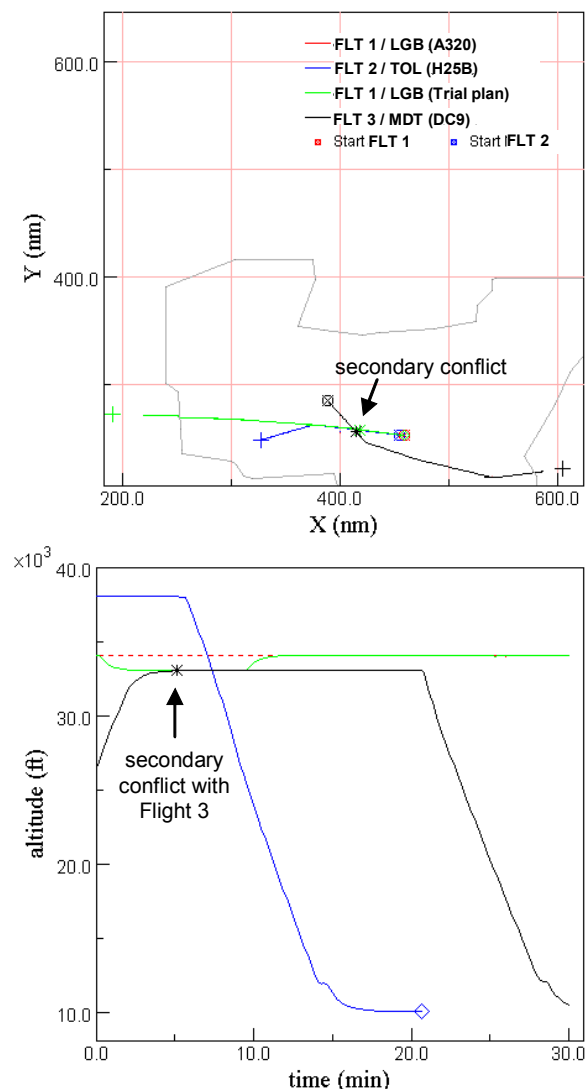


Figure 4. Trial-plan horizontal and vertical trajectories for **Flight 1** provide safe separation from **Flight 2**; however, the vertical trial-plan trajectory creates a secondary conflict with Flight 3.

Unsuccessful in maneuvering the cruise aircraft, the algorithm next attempted to maneuver the descending aircraft. A half-dozen early-descent trajectories were computed for Flight 2 and evaluated for conflict with Flight 1. Each of these trial-plan trajectories called for Flight 2 to immediately initiate a standard descent to a temporary altitude and hold that altitude until intercepting and rejoining the originally planned descent trajectory. Temporary altitudes of 33,000 ft down to 28,000 ft were attempted (in 1000-ft increments); none was successful in safely separating the two aircraft.




Figure 5. Successful resolution trajectory for **Flight 1** averts the primary and secondary conflicts with a path stretch to the right which maintains its original cruise altitude.

With its progression of vertical maneuvers complete, the algorithm next moved to horizontal maneuvers, again beginning with the preferred maneuver aircraft (i.e., the cruise aircraft), Flight 1. The algorithm began with a path-stretch trajectory turning 15 degrees to the right, yielding a 40-sec delay to the return waypoint on the aircraft's original flight plan. As illustrated in Figure 5, this trajectory was projected to successfully maintain safe separation between the two aircraft. Therefore, the resolution trajectory was accepted by the algorithm and implemented by Flight 1.

To review, two aircraft, one in cruise, the other on arrival, were projected to lose separation in less than 7 min. The conflict resolution algorithm identified an acceptable resolution trajectory, having iterated through a progression of maneuvers and maneuver aircraft. The resolution trajectory was implemented successfully. Minimum horizontal separation was 4.4 nmi, and minimum vertical separation was 4019 ft, thereby establishing safe separation between the aircraft with a 42-sec delay to Flight 1.

Case Study #2: Arrival Conflict with Arrival-Fix Restrictions and Sequencing Constraints

A more difficult resolution problem is the case of conflicting traffic on approach to a common arrival fix. The problem is more difficult in two ways: (1) the arrival fix presents a routing constraint that limits the options for resolution; and (2) it is desirable to maintain the arrival sequence, but doing so further constrains the resolution options. Case study #2 is an example of this common type of conflict resolution problem. It illustrates the method by which the algorithm accommodates arrival-fix restrictions and how arrival-sequencing constraints are incorporated into the resolution process. In this example, it is shown how the algorithm successfully spaced and sequenced a set of three aircraft on arrival to Detroit.

The conflict presented in Figure 6 is a typical case of two arrival aircraft nearing top-of-descent for approach to Detroit, a major hub airport. Flight 1 (the blue trajectory in the figure), a DC-9, is in level flight at 28,000 ft, preparing for descent. Flight 2 (red trajectory), is in trail with Flight 1 at 34,000 ft, also preparing for descent. Both aircraft are approximately 17 minutes from the arrival fix to the west when the conflict is detected. At that time, Flight 1 is projected to arrive at the fix 94 seconds before Flight 2. (This establishes the arrival sequence that the resolution algorithm will attempt to maintain.) The two aircraft are projected to lose vertical separation in 4 min, 40 sec as Flight 2 begins its arrival descent. The closest




Figure 6. Conflict on approach to a common arrival fix

point of approach is predicted to be 1.7 nmi and 461 ft.

The objective of the resolution algorithm is to resolve this conflict while complying with the arrival-fix crossing restriction and preserving the assigned arrival sequence of all aircraft bound for that arrival fix. In this case study, we focus on how the algorithm utilizes its fault-recovery mechanisms to resolve a cascading series of arrival conflicts.

Adhering to the same resolution process illustrated in Case Study #1, the algorithm categorized the conflict and began to compute and evaluate its progression of suitable trial-plan trajectories. Note that the progression, as given in Table 2, is different than in Case Study #1, because

the conflict category is different (i.e., Case Study #2 is an arrival vs. arrival conflict, whereas Case Study #1 is an overflight vs. arrival conflict).

Table 2. Progression of resolution trial plans used to deconflict Flight 1 and Flight 2

Step	Aircraft	Maneuver	Result
1	Flight 1	Increases in cruise speed	Insufficient horiz. separation
2	Flight 1	Increases in descent speed	Insufficient horiz. separation
3	Flight 2	Decreases in descent speed	Insufficient horiz. separation
4	Flight 2	Path stretches	Secondary conflict

The progression of trial plans ended in a faulted state, as no conflict-free resolution trajectories could be found. Neither speed increases to the lead aircraft nor speed decreases to the trail aircraft were projected to achieve the necessary separation. Even a 15-kt change yielded only 3.0 nmi horizontal separation. A path-stretching resolution succeeded in resolving the primary conflict. However, as shown in Figure 7, this delay maneuver also induced a new (i.e., secondary) conflict with Flight 3 (the **black** trajectory), the aircraft in sequence *behind* Flight 2 on approach to the same arrival fix. This secondary conflict was predicted to occur near the arrival fix as the trailing aircraft eventually began to catch up with the delayed aircraft preceding it, Flight 2.

This dynamic can be observed on the timeline in Figure 8. It shows the arrival aircraft in sequence from left to right. The upper timeline shows the original estimated time of arrival (ETA) of each aircraft at the arrival fix. The lower timeline shows their revised ETAs based on the trial-plan trajectory for Flight 2. As a result, it was projected that horizontal separation between it and the trailing aircraft, Flight 3, would be compromised prior to the reaching the arrival fix. Such upstream propagation of sequencing conflicts occurs frequently during an arrival rush, where traffic is converging to a common arrival fix. Because the algorithm was unable to identify a conflict-free resolution trajectory, the algorithm transitioned into fault-recovery mode.

In fault-recovery mode, the algorithm used a variant of the *de-scoping* strategy described earlier. In short, the algorithm elected to implement the resolution trajectory for the primary conflict and accept the fact that that trajectory was predicted to cause a secondary conflict later on. *How much* later on is a key parameter. In making this compromise,

the algorithm is effectively buying time, trading an urgent problem for one that is less urgent. Urgency is measured by the time to first loss of separation (t_{los}) for each conflict. If the difference between the primary t_{los} and the secondary t_{los} is large, then the compromise is considered worthwhile; if the difference between the primary and secondary t_{los} values is small, then there is limited benefit to accepting the secondary conflict. A difference of more than 4 min is currently used as the threshold for accepting a secondary conflict.

In this case, the difference between times to first loss of separation for the primary and secondary conflicts was 13 minutes, so the algorithm accepted the resolution, despite the secondary conflict.




Figure 7. Resolution of primary conflict, while accepting a deferred secondary conflict




Figure 8. Timelines show how the delay vector used to resolve the primary conflict between Flight 1 and Flight 2 impacted the spacing at the arrival fix, ultimately causing a secondary conflict with Flight 3.

Once the primary resolution was executed, the algorithm turned its attention to solving the newly created conflict between Flight 2 and Flight 3. As before, the algorithm stepped through its progression of maneuvers and applied the first-come, first-served sequencing rule to generate trial-plan resolutions. Maneuvers to speed up the lead aircraft, Flight 2, were not considered, since any increase in its speed would violate its separation with Flight 1. This is an example of how feedback from prior trial-plans can influence—and expedite—the algorithm’s progression. Since Flight 3 was next in sequence to the arrival fix, the algorithm started with a speed-reducing maneuver for this aircraft. A reduction of 10 kt CAS during the remaining cruise segment was projected to delay the flight by 16 sec and increase the separation to 6 nmi at the arrival fix, thereby resolving the conflict.

Thus, by accepting a secondary conflict in order to resolve a more urgent primary conflict, the algorithm is able to buy time to sequence arrival traffic to cross the arrival fix. It should be noted that

this strategy is consistent with a technique in common practice among controllers today. It is not unusual for a controller to deliberately place two aircraft on an intercept course in order to resolve a more urgent conflict, provided that the action will buy him/her a reasonable amount of time. The algorithm described here requires that the secondary conflict be at least four minutes farther away than the primary conflict.

Since no new upstream sequencing conflicts were generated as a result of this iteration, this concluded the resolution and sequencing process for this set of converging arrivals. The final sequence and spacing is shown in Figure 9. However, it is not unusual for the upstream propagation of sequencing conflicts to extend beyond three aircraft during heavy arrival rushes over a congested fix.

To review, three aircraft in close proximity were on arrival to a major hub airport from the east. The first two aircraft were projected to be in conflict. Unable to resolve the conflict outright, the algorithm




Figure 9. Timelines illustrate the propagation of delay up the arrival sequence.

found a compromise solution. It resolved the primary conflict by delaying the second (of the three) aircraft, and accepted a secondary conflict with the trailing aircraft as a result. The algorithm deemed the compromise acceptable, because the time to loss of separation for the secondary conflict was significantly later than the primary. The primary resolution trajectory was implemented successfully, and a resolution to the secondary conflict was identified and implemented successfully, too. Safe separation was maintained throughout. All three aircraft overflew the arrival fix as required, and the arrival sequence was preserved. The example assumed a simple first-come, first-served sequence, but the algorithm is expected to work equally well if a specific sequence were provided to it by an arrival sequencing system.

The test cases presented here were selected for illustrative purposes. The question remains as to how well the algorithm performs generally. Having described and illustrated the operation of the conflict resolution algorithm, we now turn to the simulation-based evaluation and benchmarking of the algorithm.

Approach

In order to benchmark the conflict resolution algorithm with respect to the “2x-to-3x” design standard, a simulation-based test plan was developed to stress the conflict resolution algorithm with traffic densities ranging from today’s busiest airspace to that of the envisioned future. The traffic represented the full breadth and variety of conflict situations that occur—or, it is reasoned, may occur—in real-world operations now or in the foreseeable future.

Procedure

The test procedure called for the algorithm to be evaluated in a fast-time, closed-loop air traffic environment. The Cleveland Air Route Traffic Control Center (ARTCC) airspace—the airspace with the most daily operations in the U.S. National Airspace System (NAS)—was chosen for the experiment. Cleveland Center airspace and flight operations were simulated using the Airspace Concept Evaluation System (ACES) [12]. The simulation included all types of Cleveland Center air traffic: departures, overflights, and arrivals; air carrier, business, and general aviation. Flight operations over a 24-hour period were simulated based on Cleveland Center’s radar and flight plan data, as recorded on April 21, 2005 and provided by

the FAA. Extrapolated demand sets representing twice and treble the nominal demand level were generated according to the methodology established by Paglione [13]. Flight plans were replicated and offset in time to achieve the augmented demand.

The test matrix called for the conflict resolution algorithm to be evaluated in three separate test cases of increasing traffic volume (see Table 3). Each of the three demand sets (i.e., nominal traffic (“1x”), double the traffic (“2x”), and treble the traffic (“3x”)) provided 24 continuous hours of simulated air traffic transitioning through Cleveland Center airspace. The resolution algorithm was responsible for resolving all conflicts across the Center airspace; there was no subdivision of the conflicts or traffic into sectors, and there was no human support. The simulation was executed in fast time, requiring between four and twelve hours to execute each 24-hr scenario, depending on the traffic volume. The results presented exclude arrival-vs.-arrival conflicts for flights inbound to the same arrival fix. These cases require an arrival management capability not yet implemented in the simulation test bed, which is the subject of ongoing research.

Table 3. Test matrix

Test case	Traffic demand	Conflicts detected	Conflicts resolved	Avg. delay assessed	Avg. # of trial plans
1	1x				
2	2x				
3	3x				

The resolution-initiation horizon (RIH) was set to initiate the conflict resolution process once the predicted time to loss of separation was less than 8 min. The relevant data for each detected traffic conflict inside the 8-min horizon were passed to the conflict resolution algorithm. Each resolution trajectory computed by the algorithm was transmitted to ACES, which implemented the resolution trajectory for the appropriate aircraft.

Some assumptions and simplifications were made in the modeling and execution of the experiment. In the ACES simulation environment, aircraft trajectories were entirely deterministic. Aircraft executed their planned trajectories and conflict resolution trajectories perfectly. As such, trajectory conflicts could be predicted with perfect accuracy over any time horizon, and resolution trajectories could be assured to be conflict-free. To offset the absence of uncertainty in trajectory modeling, the simulation environment provided for a user-selectable buffer distance to be added to the horizontal minimum separation parameter for the

purpose of computing an acceptable resolution. For example, resolution trajectories were required to achieve minimum horizontal separation (5 nmi) plus an additional 2 nmi.

Execution of the resolution trajectories was also simplified. Negotiation of trajectories between aircraft operators and/or the air navigation service provider (ANSP) was not modeled, and neither data-link transmission delays nor pilot-action delays were modeled. Once a resolution trajectory was determined by the automation, it was executed immediately and precisely.

Detailed trajectory data and internal algorithmic parameters were recorded for each detected conflict and each trial-plan resolution trajectory considered (i.e., rejected or accepted) by the resolution algorithm. These data were archived in a relational database for post-processing and analysis, the results of which are presented in the next section.

Results

This exploratory experiment sought to benchmark the performance of the candidate algorithm in a realistic, high-volume air traffic environment. General results are summarized in Table 4. For each simulation run (i.e., “1x,” “2x,” and “3x”), the total number of flight operations and conflicts detected is cited.⁵ The “Total conflicts” result from each simulation run is indicative of the difficulty of the problem posed to the conflict resolution algorithm. The remainder of this section presents results which benchmark the performance of the algorithm in resolving these detected conflicts.

Table 4. General results

	1x	2x	3x
Total flights	7482	15,386	24,875
Total conflicts	532	1572	3099

The conflict resolution results are categorized into a safety-related benchmark and efficiency-related benchmarks. The safety-related benchmark captures the percentage of detected conflicts that were successfully resolved in each test case. The efficiency-related benchmarks include a measure of operational efficiency (mean delay assessed per resolved conflict) as well as computational efficiency

⁵ Conflicts involving two aircraft on approach to the same arrival fix were excluded from this study and are not included in the “Total conflicts” results.

(the computation time to identify a viable resolution trajectory for a given conflict). Furthermore, it was important to determine the degree to which the conflict resolution trajectories generated by the algorithm reflected the prioritization of operationally efficient trajectories over less-efficient ones. We present the safety-related results first, followed by the efficiency-related results.

Safety-Related Results

Safety is the paramount consideration for any conflict resolution algorithm. In this study, safety was measured in binary fashion based on whether or not the conflict resolution algorithm was able to solve the detected conflict with a successful resolution trajectory. The results are presented in Table 5.

Table 5. Conflict resolution results

	1x	2x	3x
a. Total conflicts	532	1572	3099
b. Resolved on first pass	531	1569	3068
c. Resolved by deferring or de-scoping	1	1	6
d. Delegated & resolved	0	2	25
e. Unresolved	0	0	0
Resolved: Number of conflicts resolved on the first pass by the conflict resolution algorithm Resolved by deferring or de-scoping: Number of conflicts unresolved following the first pass of the conflict resolution algorithm, but successfully resolved on a subsequent pass by automatically relaxing the RIH or CFH (ref. p. 4). Delegated & resolved: Number of conflicts left unresolved by the conflict resolution algorithm, thus delegated to the conflict avoidance algorithm, and resolved by the conflict avoidance algorithm. Unresolved: Number of conflicts not resolved by the conflict resolution algorithm and not resolved by the conflict avoidance algorithm.			

As described in the Overview, the algorithm may resolve a conflict at any of several stages. The most desirable outcome is that the conflict is resolved with an efficient new trajectory on the first pass of the conflict resolution algorithm. Line b of Table 5 shows that this was the case more than 99% of the time, regardless of traffic level.

If a conflict cannot be resolved at the time it is first detected, then the algorithm resorts to *deferral* and, if necessary, *de-scoping* (ref p.4). In the simulations, a very small minority of conflicts were deferred or de-scoped, as indicated in row c of Table 5. These strategies were somewhat effective in resolving these more difficult conflict cases, albeit less so as the traffic volume increased.

In cases where the algorithm failed to compute an acceptable resolution trajectory, the conflict was *delegated* to the *conflict avoidance algorithm*. As mentioned previously, conflict avoidance algorithms deliver expediency over efficiency. Specifically, they differ from conflict resolution algorithms by their expanded maneuvering envelope (e.g., maximum allowable bank angle, climb rate, etc.) and by their simpler, fail-safe architectures. In the Automated Airspace Concept (to which the conflict resolution algorithm being benchmarked belongs), the conflict avoidance function is called the Tactical Separation Assisted Flight Environment, or TSAFE. As shown in row d of Table 5, some conflicts had to be delegated to TSAFE. It is important to note that all of these delegated conflicts were first detected with less than two minutes until the first loss of separation.⁶ While it is possible for the conflict resolution algorithm to resolve a conflict in that short timeframe, the tighter maneuvering constraints imposed on the conflict resolution algorithm (as opposed to the conflict avoidance algorithm) make it difficult. Ultimately, system architects must decide the criteria for when to invoke conflict avoidance rather than relying on conflict resolution. These results may help inform that decision.

Post-simulation analysis of the conflicts delegated to TSAFE found that all of the delegated conflicts would have been solved without violating the separation standards. This result is summarized in row e of Table 5. However, because a conflict avoidance algorithm is not implemented in the ACES simulation environment, we do not have simulation data to support this analysis. We intend to validate this result once a conflict avoidance algorithm is incorporated into ACES.

In many cases, especially at the higher traffic densities, the resolution algorithm was presented with conflicts with less than 2 min to first loss of separation (t_{los}) (see Figure 10). In a deterministic simulation environment such as the one used for this study, the majority of conflicts should have been detected at or near the 8-min range of the conflict detection algorithm, such as shown in the “1x” panel of Figure 10, where 23% of conflicts were detected 7-8 min prior to loss of separation, and another 30% were detected 6-7 min prior to loss of separation. Contrast these data with the 3x case, where 25% of conflicts were detected 1-2 min prior to first loss of separation. The frequency of “late” conflict detections presented an additional challenge to the

⁶ These cases of “late” conflict alerts are discussed in more detail later in this section.

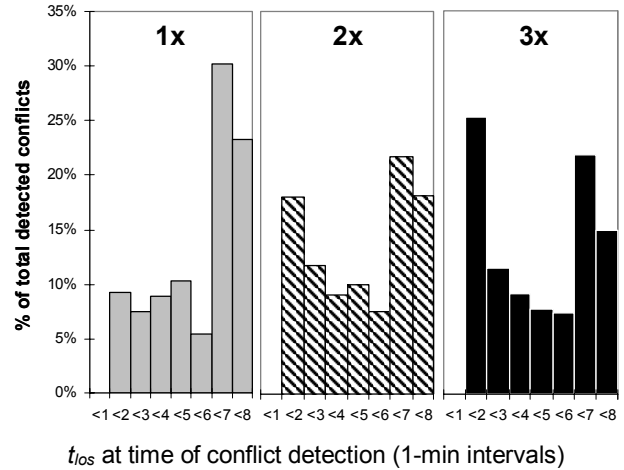


Figure 10. As traffic level increases, t_{los} at t_{detect} decreases

resolution algorithm, because near-term conflicts are more difficult to resolve than long-term ones.

The perceived “late” detections have been determined to be an artifact of the ACES simulation environment. The version of the ACES software used for this study did not perform conflict detection for aircraft outside the airspace being controlled by the algorithm (in this case, Cleveland ARTCC). As a result, many latent conflicts involving aircraft about to enter Cleveland en-route airspace were not presented to the conflict resolution algorithm until the aircraft crossed the airspace boundary, by which time the time to first loss of separation had elapsed to less than two minutes. This situation generally would not occur in a system that models handoff procedures,

Table 6. Conflict resolution results not including “late” detections

	1x	2x	3x
a. Total conflicts	490	1324	2380
b. Resolved	489	1323	2373
c. Resolved by deferring or de-scoping	1	1	2
d. Delegated & resolved	0	0	5
e. Unresolved	0	0	0

Resolved: Number of conflicts resolved on the first pass by the conflict resolution algorithm

Resolved by deferring or de-scoping: Number of conflicts unresolved following the first pass of the conflict resolution algorithm, but successfully resolved on a subsequent pass by automatically relaxing the RIH or CFH (ref. p. 4).

Delegated & resolved: Number of conflicts left unresolved by the conflict resolution algorithm, thus delegated to the conflict avoidance algorithm, and resolved by the conflict avoidance algorithm.

Unresolved: Number of conflicts not resolved by the conflict resolution algorithm and not resolved by the conflict avoidance algorithm.

where neighboring separation authorities confirm that there are no impending conflicts prior to the transfer of control. The impact of this simulation artifact was to reduce the times to first loss of separation for a significant proportion of the detected conflicts. This modeling deficiency exposed the algorithm to conflict scenarios that were more challenging to resolve than otherwise would be expected in a future operational system.

Noting that the algorithm was successful in resolving 97.5% of these “late” detections, these cases still comprised a significant proportion of the deferred, de-scoped, and delegated resolutions. Table 6 restates the results, excluding those conflicts that were detected with fewer than 2 min to first loss of separation. In the “1x” and “2x” test cases, 100% of conflicts were resolved by the conflict resolution algorithm, without need to delegate to the conflict avoidance algorithm. Delegations were reduced from 25 to 5 at the 3x level.

Efficiency-Related Results

Second in importance to safety is the operational efficiency of the resolution trajectories produced by the algorithm. Successful resolution trajectories which impose less delay or consume less fuel are preferable to those which impose or consume more.

The histogram in Figure 11 provides insight into the delay imposed by the algorithm. The histogram includes the delay data from each of the three simulation runs. The histogram is partitioned into 30-second bins (i.e., -15 sec to +15 sec of delay, +15 sec to +45 sec of delay, etc.). Negative minutes of delay indicate that the resolution trajectory realized a time savings relative to the original flight plan. The delay distributions are consistent between runs. The mean delay at each demand level is less than 30 sec and, as shown in Figure 12, appears to be stable with respect to projected increases in demand. As expected, the standard deviation of the delay histograms increases moderately with demand, from 31 sec to 39 sec to 48 sec, respectively. In the 2x and 3x cases, 21% of conflicts were resolved with a time savings (i.e., negative delay); 4% of conflicts required a delay of 2 min or more at the 3x level, 3% at the 2x level, and 1% at the 1x level.

For operational as well as computational efficiency, it is desirable that the algorithm’s progression through trial-plan resolutions reflect the operational priority of efficient solutions over less-efficient ones. In other words, since the algorithm is




Figure 11. Histogram of delay assessed by the conflict resolution algorithm

designed to select the first successful resolution trajectory it finds, it is desirable that those maneuvers near the top of the list be the most efficient ones.

Figure 13 suggests that this is indeed the case for this algorithm. The three panels correspond to the three test cases, with 1x at the left and 3x at the right. Each panel presents two types of data. The bar-chart data form a histogram (left y-axis) of delay assessed by the algorithm. This data is the same as that presented in Figure 11. Overlaid on each bar of the histogram is a diamond symbol, read from the right-hand y axis, indicating the average number of trial plans attempted for conflicts that were resolved with that amount of assessed delay. For example, in the “3x” panel at the right, conflicts resolved with zero delay (± 15 sec) required three resolution attempts, on average; but, conflicts resolved with 2 min of delay required 11 resolution attempts, on average.




Figure 12. Average delay assessed by the algorithm increases moderately with demand

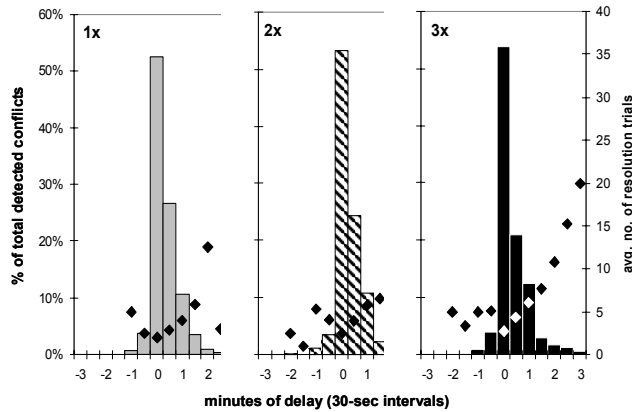


Figure 13. Histogram of delay assessed (overlay of average number of trial plans)

Figure 13 indicates that the most efficient conflict resolutions—those corresponding to the left half of each histogram—required relatively few resolution attempts, on average. Less efficient conflict resolutions—those represented at the right end of each histogram—required many more resolution attempts, on average. The correlation between more resolution trials and higher delay may reflect the additional effort typically required to avoid secondary conflicts. This phenomenon is especially evident in the 3x case, as one might expect. These data suggest that the algorithm is proposing the more-efficient resolutions first, then resorting to less-efficient resolutions only as necessary, when the more-efficient trial plans are not successful.

It should be noted that the number of resolution trials per conflict has not approached the point of becoming a limiting factor on the system. The most computationally laborious conflict in the “3x” test case required 44 trial plans. The current complement of processors being used to execute the simulation are capable of three-times this number of iterations while maintaining a better-than-real-time rate of execution.

Summary

Automated separation assurance is believed to be a key enabler in expanding the capacity of the air transportation system. System designers have established target capacity goals as high as three-times that of today. While research by the controls community in pursuit of automated conflict detection and resolution algorithms has been ample, results to date have not adequately probed the performance envelope with the volume and complexity of air traffic envisioned in the future air traffic system.

This study provides an initial benchmark for the performance envelope of one candidate conflict resolution algorithm. In fast-time simulations of the United States’ busiest en-route airspace, the algorithm successfully resolved 100.0% of detected conflicts at the 1x demand level, 99.9% at the 2x demand level, and 99.2% at the 3x demand level. Excluding conflicts that were first detected with less than 2 min to first loss of separation, which exaggerated a shortcoming of the simulation, the algorithm resolved 100.0% at 1x and 2x, and 99.8% at 3x. It must be noted that these metrics were taken in simulations having no trajectory-modeling error; to compensate, a 2-nmi buffer was applied to the horizontal separation standard for all resolution trajectories. Arrival-vs-arrival conflicts were excluded from the analysis, but will be included in a follow-on analysis.

The efficiency metrics were less sensitive to traffic level than we anticipated, an encouraging result. Average delay incurred per conflict was 21 sec at 1x traffic, rising to 25 sec at 3x traffic. Seventy-five percent of conflicts were resolved with less than a minute of delay, and an additional 21% of conflicts were resolved with negative delay (i.e., a time savings). Comparable delay distributions were found at 2x and 3x.

To provide a fair and complete set of performance benchmarks for the subject algorithm, the study should be revisited once the simulation’s modeling of handoff procedures is corrected and appropriate trajectory uncertainty models have been incorporated. At that time, the study should include analysis of arrival-vs-arrival conflict resolutions and sensitivity studies for various parameters, including the resolution-initiation time horizon and the conflict-free time horizon.

Our purpose in authoring this paper was to establish a benchmark for automated conflict resolution in a realistic, high-volume air traffic environment. It is hoped that others in the community will collaborate in this effort by documenting the performance of alternative algorithms against these same benchmarks, and by offering additional, perhaps more insightful, benchmarks for future studies here and at large. In this way, we hope to see the community begin to close in on a viable solution for automated separation assurance.

Acknowledgement

The authors wish to acknowledge the contributions of Mr. Christopher Farrell to this work.

References

- [1] FAA, 2006, *Terminal Area Forecast*, Washington, District of Columbia, USA.
- [2] FAA Joint Planning & Development Office, 2006, *Concept of Operations for the Next-Generation Air Transportation System*, Draft 3, Version 1.1a, Washington, District of Columbia, USA.
- [3] Kuchar, James K., Lee C. Yang, *A Review of Conflict Detection and Resolution Modeling Methods*, 2000, IEEE Transactions on Intelligent Transportation Systems, Vol. 1, No. 4, December 2000, pp. 179-189.
- [4] Mondoloni, Stephane, Sheila Conway, 2003, *Airborne Conflict Resolution for Flow-Restricted Transition Airspace*, AIAA-2001-4054, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.
- [5] Vivona, Robert A., David A. Karr, David A. Roscoe, 2005, *Pattern-Based Genetic Algorithm for Airborne Conflict Resolution*, AIAA-2006-6061, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.
- [6] Bilimoria, K.D., H.Q. Lee, 2002, *Aircraft Conflict Resolution with an Arrival Time Constraint*, AIAA-2002-4444, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.
- [7] Mondoloni, Stephane, Mark G. Ballin, 2001, *An Airborne Conflict Resolution Approach Using a Genetic Algorithm*, AIAA-2003-6725, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.
- [8] Erzberger, Heinz, 2004, *Transforming the NAS: The Next Generation Air Traffic Control System*, NASA TP-2004-212828, Moffett Field, California, USA.
- [9] Erzberger, Heinz, 2001, *The Automated Airspace Concept*, Proceedings of the 4th USA/Europe Air Traffic Management R&D Seminar, Santa Fe, New Mexico, USA.
- [10] Erzberger, Heinz, 2006, *Automated Conflict Resolution for Air Traffic Control*, Proceedings of the 25th International Congress of the Aeronautical Sciences (ICAS), Germany.
- [11] Erzberger, Heinz, Russell Paielli, D.R. Isaacson, and M.M. Eshow, 1997, *Conflict Detection and Resolution in the Presence of Prediction Errors*, 1st USA/Europe Air Traffic Management R&D Seminar, Eurocontrol and Federal Aviation Administration, Saclay, France.
- [12] Sweet, Douglas N., Vikram Manikonda, Jesse S. Aronson, Karlin Roth, Matthew Blake, 2002, *Fast-Time Simulation System for Analysis of Advanced Air Transportation Concepts*, AIAA-2002-4593, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.
- [13] Paglione, Mike M., Robert D. Oaks, Karl D. Bilimoria, 2003, *Methodology for Generating Conflict Scenarios by Time Shifting Recorded Traffic Data*, AIAA-2003-6724, American Institute of Aeronautics & Astronautics, Reston, Virginia, USA.